

## Introduction to Artificial Intelligence Homework 4

112550139 簡士原

### 1. Implementation Details

#### Environment:

I implemented a k-armed Gaussian bandit environment. To achieve this, I implemented 3 functions required by the PDF file:

- `reset()`: Initialize the bandit with  $\mu$  sampled from  $\mu_i \sim N(0,1)$  and clear the logs of actions and rewards.
- `step()`: Execute one pull of the selected arm and return the observed reward.
  - For the stationary case, each arm's reward is drawn from a Gaussian distribution  $N(\mu_i, 1)$ .
  - For the non-stationary cases, which means `stationary=False`, all  $\mu$  will perform a Gaussian random walk at every time step, i.e.,  $\mu_i \leftarrow \mu_i + \epsilon, \epsilon \sim N(0, 0.01^2)$ .
- `export_history()`: Return all past actions and rewards.


#### Agent:

I follow an  $\epsilon$ -greedy action-value strategy, utilizing 3 methods:


- `reset()`: Setting estimate vector  $Q$  and counter  $N$  to all 0s.
- `select_action()`: Choose an arm according to an  $\epsilon$ -greedy policy.
  - Exploration: With probability  $\epsilon$ , select a random arm.
  - Exploitation: Otherwise, choose the best action where an arm whose current estimate  $Q_i$  is maximal. If there are multiple best actions, randomly choose one.
- `update_q()`: Update the value estimate of the chosen arm,  $a$ , just pulled.
  - Learning rate  $\eta$ : If a fixed learning-rate  $\alpha$  is provided, use  $\eta = \alpha$ .  
Otherwise, use the sample-average rate:  $\eta = \frac{1}{N_a}$ .
  - Update method:  
$$Q(a) \leftarrow Q(a) + \eta \cdot (R - Q(a)).$$
  
Increment counter  $N_a$  by 1.

## 2. Experiment Results

- Part 3: Setting stationary=True, sample-average,  $\epsilon=\{0, 0.01, 0.1\}$




- A. Pure exploitation ( $\epsilon=0.0$ ) quickly plateaus on a suboptimal arm, so both its average reward and optimal-action rate stagnate at the lowest level.
  - B. Moderate exploration ( $\epsilon=0.01$ ) discovers the optimal arm more reliably than  $\epsilon=0$  while barely sacrificing early reward, leading to higher final reward and optimal-action percentage.
  - C. High exploration ( $\epsilon=0.1$ ) finds the optimal arm fastest, giving the highest optimal-action rate, and ultimately surpasses  $\epsilon=0.01$  in average reward once the cost of exploration is offset.
- Part 5: Setting stationary=False, sample-average,  $\epsilon=\{0, 0.01, 0.1\}$



- A. Pure exploitation ( $\epsilon=0$ ) cannot rediscover arms whose means have drifted, so both its average reward and optimal-action rate stay low and even decline.
- B. High exploration ( $\epsilon=0.1$ ) reacts quickly at first, achieving the highest early optimal-action rate, but its constant probing adds noise. Combined with the slow-adapting sample-average estimate, this causes a gradual drop in optimal-action rate and leaves its final reward below that of  $\epsilon=0.01$ .
- C. Moderate rate ( $\epsilon=0.01$ ) strikes the best balance: it refreshes knowledge often enough to track the drifting optimum, yet avoids the excessive variance of  $\epsilon=0.1$ , ending with the highest average reward and the most sustained


optimal-action percentage.

- Part 7: Setting stationary=False,  $\alpha=0.1$ ,  $\epsilon=\{0, 0.01, 0.1\}$



- Pure exploitation ( $\epsilon = 0$ ): With  $\alpha$  fixed at 0.1, every new reward quickly influences the value estimate, so even  $\epsilon = 0$  can climb well above its Part-5 performance; but without exploration, it levels off lowest in both metrics.
  - High exploration ( $\epsilon = 0.1$ ) keeps finding the current best arm fastest, giving the top optimal-action percentage ( $\approx 0.80$ ). However, its frequent random pulls add variance, so its average reward lags slightly behind  $\epsilon = 0.01$ .
  - Moderate exploration ( $\epsilon=0.01$ ) strikes the best balance: enough exploration to track drifting optima, yet few enough random pulls that its average reward ends highest while its optimal-action rate stays close to that of  $\epsilon = 0.1$ .
3. Discussion:

I ran one more little test just accidentally: kept  $\epsilon = 0$  and only changed the step-size  $\alpha$ . Turns out  $\alpha = 0.1$  outperforms  $\alpha = 0.01$  significantly.



I think it's all about how fast the Q-values can 'forget' old rewards: with  $\alpha = 0.1$ , the agent's memory only stretches back about 10 pulls, so it notices the drift almost right away and jumps to the new best arm. In contrast,  $\alpha = 0.01$  clings to a 100-step history, basically sleeping through the change and sticking to a now-suboptimal arm for ages.